

e-YSIP 2022

# FPGA FOR EDGE



Anupam Kurien Mathew  
Dan Mani Binu  
Hari Vikinesh

## **Mentors**

Isha Kamone  
Lohit Penubaku

Duration of Internship: 06/06/2022 – 23/07/2022

*2022, e-Yantra Publication*

# FPGA for Edge

## Abstract

FPGAs are steadily entering into the race of compute-intensive work precisely in the field of Machine Learning which is referred to as hardware accelerators. These accelerators can improve the performance of the ML Model significantly compared with traditional computing devices like CPU and GPU. FPGAs are preferred because of the hardware reconfigurability, power efficiency and low cost compared to CPUs and GPUs. This project aims to understand and exploit the potential of FPGAs in the application of agriculture. The project focused on training ML models which can be deployed on FPGA, and video processing to use the ML models for indoor agricultural applications. At First, the performance of various ML Models is studied for agricultural applications. In this project ML model is studied and developed for tomato leaf detection. Though there are multiple ML Models available on the air to use, none of them can be used directly on the FPGAs. The ML Model to be chosen should meet the resource constraints of the FPGAs that are bound to be used. Also, FPGAs require sophisticated architecture to run any ML Model. The development of application-specific architecture improves the performance of the model significantly. So, it is important to choose an appropriate ML model to use on FPGA. By analyzing various algorithms considering the constraints. In this project, a YOLOv2-Tiny model is developed from scratch to detect tomato leaves using FPGA and their performance parameters are being measured. A comparison study between FPGA, CPU & GPUs is also added to this project. This comparison study shows the performance improvement that the FPGA can afford.

## Completion status

Though, this project aimed to do video processing on FPGA. Unfortunately, this project ended up deploying ML models on images and comparing the results of it with CPUs and GPUs.



## 1.1. HARDWARE PARTS

### 1.1 Hardware parts

This section discusses the hardware parts that are being used in this project along with their descriptions and key features. In this project, two FPGA SoC boards are utilized to measure the performance of the ML Model.

#### ZedBoard

ZedBoard is one of the famous FPGA SoC released on 2011 by Avnet. It has an ARM processing system along with Zynq-7000. FPGA Part – XC7Z020-clg-484-1. Features of ZedBoard are as follows.

- 512MB DDR3
- 256Mb Quad-SPI Flash
- Onboard USB-JTAG Programming
- 10/100/1000 Ethernet
- USB OTG 2.0 and USB UART
- PS & PL I/O Expansion(FMC, 5\*Pmod, XADC)
- Displays(1080p HDMI, 8-bit VGA, 128x32 OLED)
- I2S Audio CODEC

A detailed description of this board is given in the reference [manual](#) and this board can be purchased using this vendor [link](#). Figure 1.1 shows zedboard.




Figure 1.1: ZedBoard

## 1.1. HARDWARE PARTS

### Nexys Video board

Nexys Video is based on Artix-7 FPGA released on 2015 by Digilent. A softcore Microblaze processor is present as processing system integrator. FPGA Part – XC7A200T-1SBG484C. Features of Nexys Video Board are as follows.

- 512MB DDR3
- 32Mb Quad-SPI Flash
- USB Host, USB UART and USB programming
- 10/100/1000 Ethernet
- 3.75Gbps GTP transceivers
- FMC, 4\*PMOD & XADC
- Displays(1080p HDMI In, HDMI Out, 128x32 OLED)
- I2S Audio CODEC

A detailed description of this board is given in the reference [manual](#) and this board can be purchased using this vendor [link](#). Figure 1.2 shows Nexys Video Board.




Figure 1.2: Nexys Video Board



### 1.2 Software Setup

This section describes all the software that is required to run the project on the FPGA and measure the output performance of the CPU & GPUs.

- **Vivado HLS** – This software is used to write C/C++ codes in it and then it can synthesize the code to RTL design. It is also possible to simulate the C Code using a C simulator as well as RTL Simulator. After simulation and synthesis, this software can convert the code into an IP block.
- **Vivado** – This software acts as a top entity and this software acts as an architect to design the top-level module design and generate bitstream to implement them on hardware. In this software, the generated IP blocks can be used to create a design or users can use HDL languages like Verilog or VHDL to generate a bitstream file. One main advantage of this software is it supports mixed language implementations which means users can use both Verilog and VHDL in the same design. The top-level design can be exported as hardware design as well.
- **Xilinx SDK** – This software is used to run the hardware design as a bare-metal application on the FPGAs. This software is used if the user has processing system parts involved in the design. The user should create a separate project to program this processing system part.
- **Paperspace** – Paperspace is a cloud computing service developed for Machine Learning and other compute-intensive tasks to run on the cloud GPUs. This organization provides Cloud GPUs at a low cost. They provide CPUs and GPUs for usage at cost on an hourly/monthly basis. The GPUs are ranges from Nvidia M4000 to A100 models. In this project, an Nvidia A4000 model is provided for the development and testing of the ML Model.

Apart from the software mentioned above. There are certain python modules needed to be installed to run the model which are mentioned below.

- matplotlib - interactive graph creator module
- Tensorflow - ML Traing and Testing module
- pyserial - Serial Read module

Also, powerstat linux tools are required to measure the power consumption.



## 1.2. SOFTWARE SETUP

### Installation Steps

1. Visit this [link](#). Scroll down and click on 2018.3 Archive.
2. Choose the third option to download the software. Refer Figure 1.3.

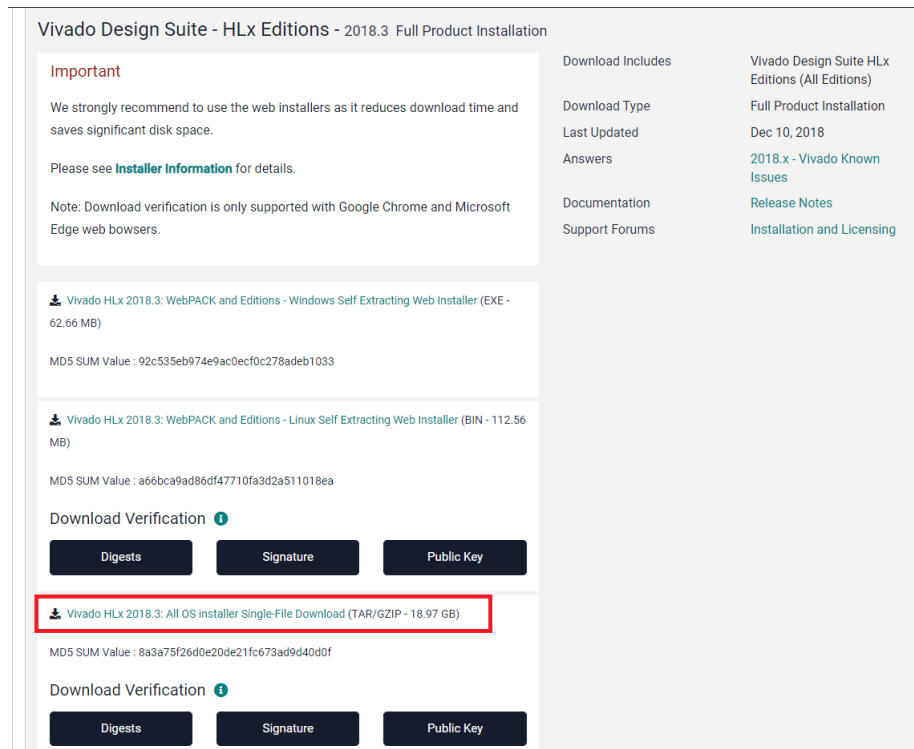


Figure 1.3: Xilinx Vivado Download Page

3. If a login prompt appears log in using the Xilinx account and the download will start automatically.
4. After downloading the software, Extract it. Now, the xsetup.exe file should be present inside the folder.
5. After the installer appears. On the First Page, Click on Next. On the license agreement tab, accept all license agreements and click on Next.
6. Now, Choose Vivado HL Design Edition and click on Next.
7. Click on Next and in the next tab choose the appropriate location (prefer SSD) to install the software.
8. Finally, on the next tab click on the Install button to install the software.

## 1.3 Model Development

This section of the report explains the development of the custom model for tomato leaf detection using object detection algorithms. The details for building a Hardware accelerator are also described in detail.

### Object Detection

Object detection is a computer vision technique for locating instances of objects in images or videos. All the Object detection algorithm uses Machine Learning Techniques to detect some objects and produce meaningful results. Some popular object detection algorithms are R-CNN, SSD, and YOLO algorithms. To detect an object using these algorithms requires a pre-trained network. This pre-trained network is called a base model or weights of the network. In this project, to detect tomato leaves YOLOv2-Tiny model has been used. At the start of the project, different tomato leaf detection models are developed such as YOLO-v5, YOLO-v3 tiny and mask-r-CNN. The results of these models are shown in the figure 1.4. At Later stage of the project models for yolo-v2 and SSD mobilenet are also developed. In comparing the results from various models with the model size, the YOLOv2-Tiny model is chosen. The model size is a critical parameter because FPGAs have limited resources.




Figure 1.4: Initial Model Results

### 1.3. MODEL DEVELOPMENT

#### YOLOv2-Tiny Architecture

YOLOv2-Tiny gives state-of-the-art detection accuracy on the PASCAL VOC and COCO. It can run on varying sizes offering a tradeoff between speed and accuracy. The dataset used for this project is the village plant dataset and modified to PASCAL VOC format only for healthy tomato leaves. This model is trained in the Paperspace platform. The figure 1.5 shows the architecture of YOLOv2 that is used in this project.

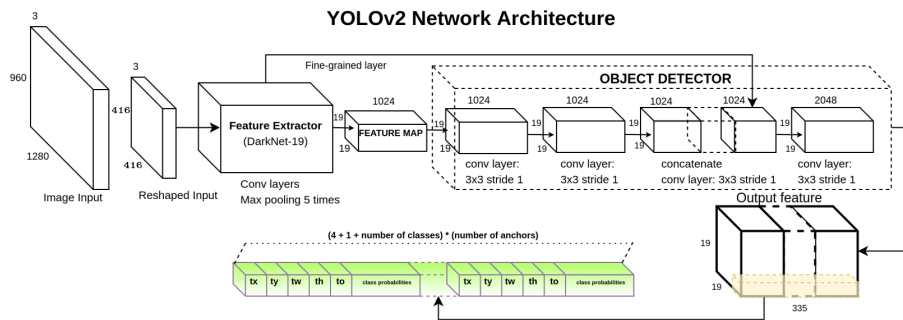


Figure 1.5: YOLO-v2 Architecture

#### Hardware Accelerator

An Object detection model cannot be implemented on the FPGAs directly. So, a Hardware Accelerator must be developed to run the model. So, a custom hardware accelerator for FPGA is developed to detect tomato leaves using the YOLOv2-Tiny model. This Accelerator uses Vivado HLS pipelining tools to reduce the overall execution time significantly and manages memory utilization. The basic architecture of our model is shown in the figure 1.6. The output result of this accelerator is sent to the laptop for visualization using UART Communication.




Figure 1.6: Hardware Accelerator





### 1.4 Code Explanation

All the codes that are used in this project are described in this section. The code explanation part is split into four parts namely python notebooks, python scripts, HLS Codes and Verilog HDL. Python scripts are used to test the model and obtain the results of the YOLO-v2 model. HLS Codes are used to develop the hardware accelerators for FPGA deployment. Verilog HDL is used to transmit the output of the hardware accelerator to the laptop. Necessary files and codes to replicate this project is added to the GitHub repository and it can be accessed using this [link](#).

#### Python Notebooks

##### Notebook 1 - yolov2-tiny.ipynb

This Notebook is used to train the yolov2 algorithm using the darknet. This notebook has the necessary comments to develop a custom model and convert it to the Keras model at the end. This notebook uses the pre-trained network to train a new model. Important code snippets for training and testing is mentioned below.

```
# Start Training
# users are free to use custom configurations or use
# prebuilt configuration files under /darknet/cfg
# directory

!./darknet detector train yolov2.data \
    cfg/<custom_cfg>.cfg \
    yolov2.weights -dont_show

# Test the trained model using a darknet detector
# and use appropriate paths
!./darknet detector test yolov2_tiny.data \
    cfg/yolov2_tiny.cfg \
    backup/yolov2_tiny_last.weights \
    test/11.jpg

# Convert weights to a Frozen Graph
# Darkflow repository is used
!./flow --model /<path>/yolov2_tiny.cfg \
    --load /<path>/yolov2_tiny.weights --savepb
```



## 1.4. CODE EXPLANATION

### Notebook 2 - yolov2-scratch.ipynb

This notebook is used to develop and test the yolov2 algorithm from scratch as it is using very minimum external modules. Also, the functions for the algorithm are written in such a way that they can be converted to C/C++ code for HLS implementation. This notebook has all the functions to load the yolov2 model and detect bounding boxes using the yolov2 regressor.

```
# Recreate the exact same model, including its weights
and the optimizer
new_model = tf.keras.models.load_model('yolov2-tiny.h5')

# Show the model architecture
new_model.summary()

def padding(X, pad, n_H, n_W, n_C):
    the function creates a layer of zeroes around
    the entire 2D matrix. The following function
    is modified for padding in 1 dimension.
    X — array to be padded with zeros
    pad — the width of padding
    n_H — height of the 1D array in 2D form
    n_W — width of the 1D array in 2D form
    n_C — no of channels in image
    X_pad — output array which whose dimensions
    have been increased to (n_H/n_C + 2*pad)

def convolution(k_width, k_height, filters, arr_width,
               arr_height, arr_channels, arr, s, k):
    performs normal convolution on the array/image passed
    the function performs a horizontal convolution first
    through the entire image and then performs vertical
    convolution
    k_width — kernel width
    k_height — kernel height
    filters — the no of kernel filters applied
    to the image/array
    arr_width — width of the array
    arr_height — height of the array
    arr_channels — no. of channels in the passed array
    arr — the image/array on which convolution
```



## 1.4. CODE EXPLANATION

is to be performed  
s                   – stride  
k                   – the kernel weights as 1D array  
arra               – final output after convolution

```
def batch_norm(X, arr_channels , arr_height , arr_width ,K):
```

batch normalization performs –

1. subtraction with moving mean indicated by  $K[2*arr\_channels+j]$
2. division with root of standard deviation indicated by  $math.sqrt(K[3*arr\_channels+j])$
3. multiplication by Gamma indicated by  $K[j]$
4. addition by Beta indicated by  $K[1*arr\_channels+j]$

X                   – array to be batch normalized  
arr\_channels       – no.of channels in array  
arr\_height         – height of the array in 2D  
arr\_width          – height of the array in 2D  
K                   – Batch normalization weights in the order gamma, beta, standard deviation, subtraction  
X                   – calculations performed in the same array and returned as output

```
def max_pooling(k_width = 2, k_height = 2, filters = 3,  
arr_width = 8, arr_height =8 ,arr_channels = 3,  
arr = img,s = 2):
```

Max Pooling calculates the maximum value for patches of a feature map, and uses it to create a downsampled feature map. In and Outs are similar to convolution.

```
def leaky_relu(alp, arr):
```

leaky relu returns same number if the number is greater than 0 else returns the number after multiplying it alp.  
alp – computation parameter  
arr – the array on which leaky relu is to be performed  
arr – also returned as output



## 1.4. CODE EXPLANATION

### Python Scripts

#### Python Script 1 - visualize.py

This code is used to get the results from the Hardware accelerator of the FPGA using the serial communication protocol(UART). This code reads the data from the serial port and decodes it to show the detection result using OpenCV. Main parts of the code are attached below.

```
# Open Serial Port
ser = serial.Serial('COM6', 115200, timeout=1)
# Reading Serial data
while True:
    # Reading a Line Input
    message = ser.readline()
    if message:
        # Converting Byte String to unicode string
        data = message.decode()
        break

# Draw Bounding Box and label it
img = cv2.rectangle(img, (int(data[3]), int(data[4])),
                    (int(data[5]), int(data[6])), (255,0,0), 1)
img = cv2.putText(img, data[2], (int(data[3]),
int(data[4])+25), cv2.FONT_HERSHEY_SIMPLEX,
0.9, (0,0,255), 1)
```

#### Python Script 2 - convertImage.py

This code is used to convert the input RGB image into a single dimension array. This array is stored in the C header file using this script.

```
# Function File Header Writer
def file_write(file, data, h, w, c):
    This function will write a header file to store images.
    file - path with filename
    data - image data
    norm - Normalization
    h - height
    w - width

# Converting 3d Array to 1d list
for i in range(0, h):
    img_linear = []
```



## 1.4. CODE EXPLANATION

---

```
for j in range(0, w):
    for k in range(0, c):
        img_linear.append(img[i][j][k])
row_array.append(img_linear)
```

### Python Script 3 - yolov2-tiny.py

This code is used to run the yolov2-tiny algorithm on the CPU and GPU to analyze the performance of the model. The yolov2-tiny model is converted to TensorFlow frozen graphs for implementation. The output of this code returns the leaf detection bounding boxes along with the accuracy and latency.

```
def main():
    This function loads the yolov2-tiny model using the
    tensorflow. After that, this function will call
    other funtions to detect the tomatoleaf in given
    image. The Output image is displayed using the
    matplotlib library. This function also calculates
    the latency of the model.

def get_input_data(img_file):
    This function will read the input image using the
    PIL library and convert the image to a normalized
    float32 numpy array.
    input:
        img_file - path to input image
    returns:
        new_image - resized image
        input_data - normalized numpy array

def check_result(data):
    This function interprets the numpy array and apply
    bounding box regressor of the yolo architecture to
    finalize outputs. This function instatiates small
    math functions to calculate results.
    input:
        data - normalized numpy array
    returns:
        output - predictions
```



## 1.4. CODE EXPLANATION

### HLS Codes

#### HLS Code 1 - yolov2.cpp

This code is the brain of the custom hardware accelerator and this code detects the tomato leaf based on the trained model and input image using the yolov2 architecture. The regressor for yolov2 algorithm is described below.

```
// Defining Custom Datatype
typedef ap_int<9> bbxy;
typedef ap_int<10> accuracy;
typedef ap_int<1> int1;
typedef ap_int<3> int3;

// Declaring Output Products
bbxy bbx1, bby1, bbx2, bby2;
accuracy score = 0;
int3 leaves = 0, flag = 0;
int1 result = 0;

void yolo_eval(float *box_xy, float *box_wh,
               float *box_confidence, float *box_class_probs,
               float *boxes, float *box_class_scores,
               float threshold, int &count)

    Function Name: yolo_eval
    Input      :
    — Yolo Outputs —
    box_xy      — bounding box x and y position
    box_wh      — bounding box width and height values
    box_confidence — confidence of the bounding boxes
    box_class_prob — detected class probability
    —————
    boxes        — maximum predicted box-whs
    box_class_scores — bounding box class scores
    threshold    — threshold for scores
    count        — counter
    Description  : This function converts yolo outputs
    (multiple bounding boxes) to best predicted bounding
    box along with their scores, box co-ordinates and
    classes.
```



#### 1.4. CODE EXPLANATION

---

```
void non_max_suppresion(float *scores, float *boxes,
    int count, int max_boxes = 10,
    float iou_threshold = 0.3)
```

Function Name: non\_max\_suppresion

Input :

scores — scores array

boxes — bounding boxes array

count — count value

max\_boxes — maximum no of bounding boxes

iou\_threshold — threshold for intersection over union

Description : This function validates all the input bounding boxes and outputs the more appropriate bounding box for the input parameters.

```
int yolov2Algorithm()
```

Function Name: yolov2algorithm

return : output flag

Description : Function to perform yolo operations it calls appropriate functions to do convolution for the given input images.

```
int regressor()
```

Function Name: regressor

return : output flag

Description : This function stores parameters for yolov2 regressor and calls appropriate functions to complete evaluation and give the final results.

```
void hlsTop(int1 start, int1 &compute, int3 &leaves,
    accuracy &scores_out, bbxy &out_bbx1, bbxy &out_bby1,
    bbxy &out_bbx2, bbxy &out_bby2)
```

Function Name: hlsTop

Input : start — start switch

Output :

compute — start tx computation flag

leaves — no of leaves detected flag

scores\_out — accuracy of detected leaves

out\_bb— — bounding box co-ordinates

Description : This function will create the module I/Os using the HLS Pragmas.



## 1.4. CODE EXPLANATION

### Verilog HDL

#### Verilog Code 1 - uart\_tx.v

This module is responsible to get input from the ssd model and convert the decimal value to digits for UART transmission. Then, this module transmits the necessary signals for the uart\_tx module to send the data serially.

##### Inputs

clock           – 100Mhz Clock  
reset           – restart computation  
result          – model identification flag  
compute        – start flag for computation  
imgID          – image Number  
accuarcy        – Model's score  
bbxn, bbyn      – co-ordinates of the boundary boxes

##### Outputs

tx              – serial data output pin  
txComplete     – transmission status

#### Verilog Code 2 - uart.v

This module is responsible for transmitting the message through serial communication protocol UART.

##### Inputs

clk            – 100MHz on-board Oscillator  
reset          – reset to resend serial data  
result         – tells the tomato leaf detected or not  
imageID        – no of the image used  
acc\_n          – accuarcy in three digit  
bb1\_n          – boundary box 1 co-ordinates in six digits  
bb2\_n          – boundary box 2 co-ordinates in six digits

##### Outputs

txOut          – Serial Output  
txDone         – Transmission Status

##### UART Transmission Parameters

baud rate = 115200  
clocks per bit (cpb) =  $(100 \times 10^6) / 115200 = 868$   
start bit = 0, stop bit = 1;  
1 start bit, 1 stop bit and no parity bit  
// msg Format = #2-1-97.8-213-231-345-322#  
– is the delimiter  
# is start and stop bit indexes





## 1.5 Demonstration

This section deals with implementing the hardware accelerator on the FPGA using the step-by-step method. Also, this section lists the steps to obtain the results from the CPU and GPU for comparing the results. A demonstration video is also added to Google [Drive](#) for the users to get a better understanding of the implementation process. All the project files, documents, reference materials and models that are used and carried out throughout the project are added to this Google [Drive](#).

### Deploying the Hardware Accelerator

After generating bitstream using the vivado software, it is ready for deployment. But, before uploading bitstream into the FPGA. The users should run the python program on their laptop or PC. The steps for executing this program are mentioned below with the reference figures.

1. Connect the micro-USB data cable to the UART port of Nexys Board and Connect a male to female jumper wire to the JA1 PMOD connector and the RX port of any USB to UART converter.
2. Open Device Manager on the PC or Laptop and check the COM port. Open the visualize.py program file and make sure that both COM ports are the same.
3. Run the Program in the command prompt or any python IDE.
4. Now, Program the FPGA with bitstream using Vivado Software.
5. After Accelerator completes the detection, the results are transmitted to the PC and the user should see an output similar to the figure 1.7.

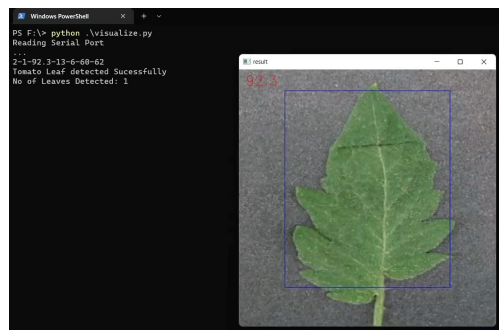


Figure 1.7: Output of the Accelerator



## 1.5. DEMONSTRATION

### Obtaining Results from the Hardware

To analyze the results with the Hardware Accelerator, the users can follow the steps mentioned below to obtain the results from the CPU and GPU.

To measure the performance of the built YOLOv2 model, the required modules to run and test the program should be installed on the computer. All the required modules are mentioned in [section 1.2](#).

To measure the performance of the model follows the steps mentioned below.

1. Make sure that tfLite model, test images and yolov2-tiny.py are in the same directory.
2. Open two terminals in separate windows, first run the python program in one terminal and the powerstat tool in another window.
3. Python program shows the results of the model in the terminal that includes accuracy, class, bounding box and latency.
4. Parallely, the power stat tool shows the power consumption of the computer, when the python program is running at the backend the power consumption should be recorded. Refer figure 1.8.

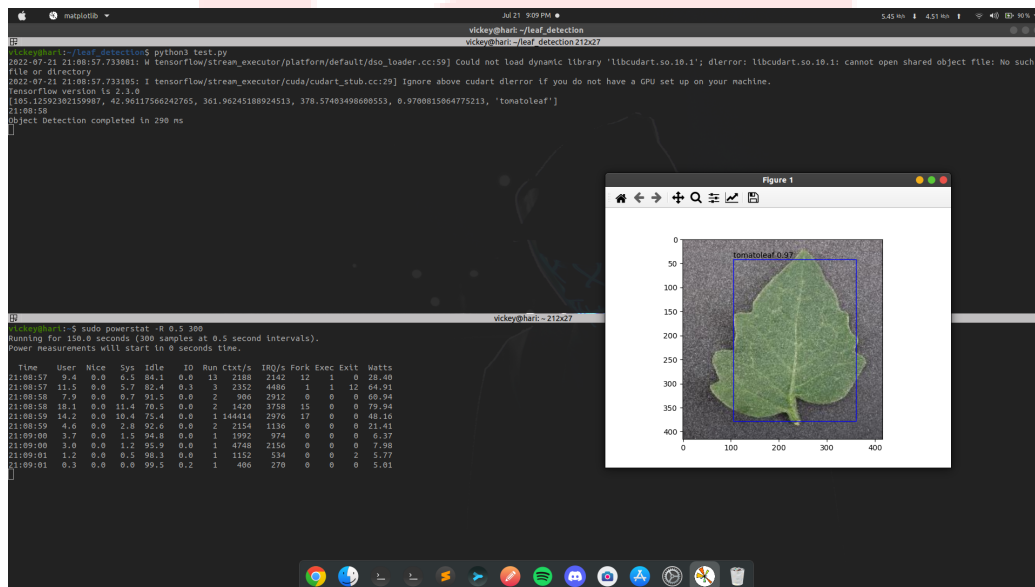


Figure 1.8: Sample Recording

## 1.6. PERFORMANCE STUDY

To measure the performance of the model on the FPGA with resource utilization Vivado softwares are used. If the Implementation of the design in the Vivado software completes successfully, then Vivado provides the complete report of the design with the details of resource utilization, peak power consumption and timing constraints. Refer the figure

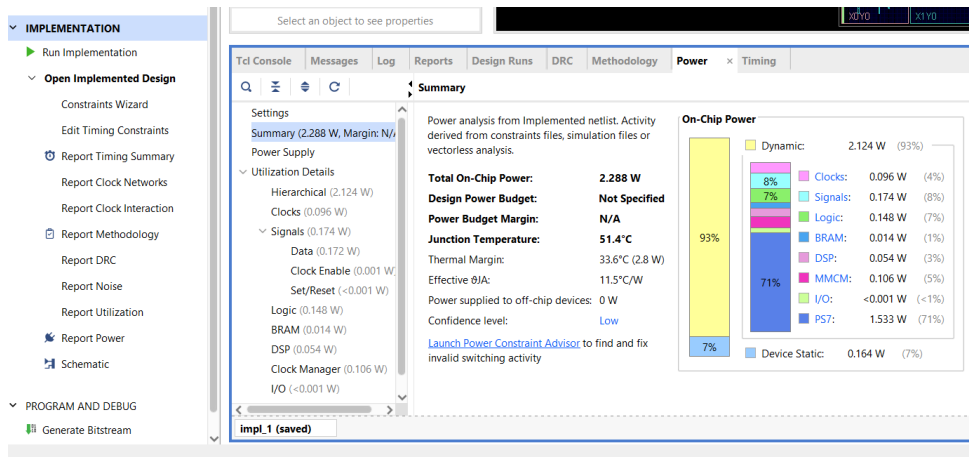


Figure 1.9: Recording Resource Utilization

## 1.6 Performance Study

This section studies the results from the implemented Hardware Accelerator and compares the performance with the CPU and GPU. Also, a performance comparison between the implemented design and the existing designs is also analyzed.

### Comparison with Traditional Hardwares

The project laid focus on the comparison between eight hardware, comprising four CPUs two GPUs and two FPGAs. Latency and Power consumption are the two domains under comparison.

From the Figure 1.10 Among the CPUs, the i7-9750CPU had the least latency of 148ms in implementing YoloV2-tiny while that of INTEL XENON 2.3GHz was recorded to be the highest (326ms). Both the ZedBoard and Nexys Video latency was recorded to be about 210ms with Nexys Video Board having better performance than Of all devices, GPUs are found to have the least latency, with NVIDIA RTX A4000 running at 124ms latency and NVIDIA Tesla at 84ms.

## 1.6. PERFORMANCE STUDY




Figure 1.10: Latency Comparison

Looking into power consumption from figure 1.11, The i7-9750H drinks about 79W of power and the i5-9400H at 59W. Both the NVIDIA Tesla K80 and NVIDIA RTX A4000 consume a reasonable amount of 43W, while the Nexys and ZedBoards consume very less power of 3W.




Figure 1.11: Power Consumption Comparison

## Comparison with Existing Accelerators

This work used 32-bit floating point representations for weights leading to achieving higher accuracy. This work requires the lowest number of flip flops at 28.5K and 23.8K for ZedBoard and Nexys video board respectively. From the table, BRAM utilization on this work is also lower than that of Kintex Ultrascale (1814 KB) at 846KB for Zedboard and 890KB for Nexys while



## 1.7. CONCLUSION

possessing a bigger image size and cnn size. The DSP utilization for Zedboard is lowest among all hardware accelerators compared at 121 while Nexys board matches that of Zynq ZC702 at 140. The hardware utilization is lower only on accelerators which use a combination of smaller image size and lower precision. In that case, the utilization is lower than Kintex Ultrascale. The other works which use higher image size consequently use significantly higher resources. This translates to those accelerators not being implementable on lower-tier FPGAs.

Table 1.1: Comparison with Hardware Accelerators

Platform	Zynq XC7Z045	Xilinx Zynq ZC702	Kintex Ultrascale XCKU115	Stratix-V	Virtex-7 VC707	Zynq Ultrascale+	Intel Arria 10 GX115	ZedBoard (This work)	Nexys Video (This work)
Frequency (MHz)	200	100	125	150	200	200	200	100	100
BRAMs (KB)	186	630	1814	2210	1214	1824	2232	846	890
DSPs	-	140	-	384	272	2520	1518	121	148
LUTs	46.3K	36.1k	392.9K	230.9K	104.7K	600K	138K	51.4K	40.2K
FFs	-	36.8K	348K	350K	140.1K	-	823.4K	31.4K	31.4K
CNN Size	0.1125	14.5	1.2	1.45	30.74	5 layers of VGG	30.95	42.3	42.3
Precision	8-bit fixed	8-bit fixed	8-bit fixed	8-bit fixed	(8-16)bit fixed	16-bit fixed	16-bit fixed	32-bit float	32-bit float
Image Size	32x32	-	32x32	224x224	224x224	224x224	224x224	64x64	64x64

## 1.7 Conclusion

From the comparisons given above, it is evident that ML model implementation works best on FPGAs and GPUs. The choice between FPGA and GPU is application specific. FPGAs are best for low-power implementation and standalone applications. GPUs for low latency models. A more efficient code can alter the speed and resources required to implement the model. Traditional hardware like CPUs has turned out to be poor choices for model implementation. Based on the resource cost and simple architecture the propped model can be implemented on low-tier FPGAs.

## 1.8 Future Work

This project can act as the baseline for the projects implementing real-time leaf detection using Xilinx FPGAs. This project's code can be implemented on Zynq and MicroBlaze series FPGAs for performance analysis. It is also possible to develop any custom object detection model and implement it on the other FPGAs using this project.



### 1.9 Bug report and Challenges

In HLS codes there are some parts of the code that are not optimized to the best. Also, the use of variables and their size can be reduced to some extent to reduce bottleneck memory usage on the FPGAs. The UART Verilog code has a decimal to digit converter that has division and modulus operations costing higher resource utilization. So, other methods should be employed to reduce resource utilization.

An Important challenge our project has is integrating the OV7670 camera into the FPGA. This camera was used at the start of the project. Though the camera interface is proper, the output from the camera is not clear. This is due to improper I2C configuration information provided by the vendor. Some FPGA hobbyists have claimed this issue on the internet. So, this area will be quite tricky for the developers.

Another challenge is trying to figure out the appropriate model to deploy on the FPGA. This is because of the resource constraints of the FPGAs. Also, memory management is really important to reduce latency for reading and writing operations. If the FPGA tries to access memory from the external DDR or other secondary memory storage devices, it will increase the latency of the entire model.

# Bibliography

- [1] S Zhang, J Cao and Q Zhang, *An FPGA-Based Reconfigurable CNN Accelerator for YOLO*, 2020.
- [2] DT Nguyen, TN Nguyen Hyuk Lee and H Kim, *A High-Throughput and Power-Efficient FPGA Implementation of YOLO CNN for Object Detection*, 2019.
- [3] Coleman, Steven, and Marian Verhelst, *High-Utilization, High-Flexibility Depth-First CNN Coprocessor for Image Pixel Processing on FPGA*, 2021.
- [4] Sharma, Hardik and Mahajan, *From High-Level Deep Neural Models to FPGAs*, 2016.
- [5] K Kara, D Alistarh and G Alonso, *FPGA-accelerated Dense Linear Machine Learning A precision Convergence Trade-off*, 2017.
- [6] MA Momen and MAS Khalid, *FPGA-Based Acceleration of Expectation Maximization Algorithm Using High-Level Synthesis*, 2019.
- [7] A Sanaullah, C Yang and Y Alexeev, *Real-time data analysis for medical diagnosis using FPGA-accelerated neural networks*, 2018.
- [8] S Jiang, D He and C Yang, *Accelerating Mobile Applications at the Network Edge with Software-Programmable FPGAs*, 2018.
- [9] S Kala, J Mathew and BR Jose, *UniWiG Unified Winograd-GEMM Architecture for accelerating CNN on FPGAs*, 2019.
- [10] X Liu, HA Ounifi and A Gherbi, *A Hybrid GPU-FPGA-based Computing Platform for Machine Learning*, 2018.